

# Comprehensive Guide to the `head` and `tail` Commands in Unix/Linux

`head` and `tail` are two of the most fundamental and frequently used Unix text-processing utilities. They are perfect companions: `head` **shows the beginning of a file/stream**, while `tail` **shows the end**. Together, they allow quick inspection of files without opening them fully — ideal for logs, large datasets, CSVs, or any text output.

This guide covers **every aspect** of both commands: full syntax, all options (GNU coreutils vs BSD/macOS differences), byte vs line mode, advanced features like live following, real-world patterns, pipelines, pitfalls, best practices, and performance notes. Everything is based on modern systems (2026 era): GNU coreutils (Linux) and BSD variants (macOS/FreeBSD).

We'll treat them together since they share many concepts (just mirrored behavior).

---

## 1. Basic Purpose & Syntax

`head` — Output the first part of filesstreams.

`tail` — Output the last part of filesstreams.

**General Syntax** (identical structure):

Bash

```
head [OPTION]... [FILE]...
tail [OPTION]... [FILE]...
```

- If no FILE (or `-` as filename), reads from `stdin` → perfect for pipelines.
- Multiple files → processes each separately, with headers by default (`-v` behavior).
- Output goes to `stdout`.

**Default behavior** (no options):

- `head` → first 10 lines
  - `tail` → last 10 lines
-

## 2. Core Options — Lines vs Bytes

Both support **line mode** (default) and **byte mode**.

Applies				
Option	Long Form	To	Meaning	Example
-n N	--lines=N	Both	N lines (head: first; tail: last)	head -n 5 → first 5 lines
-n +N	(tail only)	tail	From line N to end (skip first N-1 lines)	tail -n +100 → lines 100-end
-c N	--bytes=N	Both	N bytes (head: first; tail: last)	head -c 100 → first 100 bytes
-c +N	(tail only, GNU)	tail	From byte N to end	tail -c +1024 → after first KB

### Notes:

- N can be suffixed: K (kilobytes), M (mebibytes), G (gibibytes) — e.g., -c 10M .
- Negative N (e.g., -n -5 ) is **GNU extension** for "all but last/first N" (see advanced).

Example file log.txt (20 lines):

```
Bash
head -n 3 log.txt      # lines 1-3
tail -n 3 log.txt      # lines 18-20
tail -n +18 log.txt    # lines 18-20 (alternative)
head -c 50 log.txt     # first 50 bytes
```

## 3. Verbose / Quiet Mode (Multiple Files)

When processing **multiple files**, behavior changes:

Option	Long Form	Meaning	Default Behavior
-v --verbose	Always print filename headers	No headers (single file)	
-q --quiet / --silent	Never print headers	Headers for multiple files (default)	

Bash

```
head -v file1 file2 file3  # => ==> file1 <== \n content... \n ==> file2 <==  
head -q file1 file2      # no headers
```

## 4. Advanced tail Features (Live Monitoring / Following)

tail's most powerful feature: **real-time following** with `-f`.

Option	Long Form	Meaning	Notes
-f	--follow	Follow file as it grows (live tail)	Ctrl+C to stop; GNU & BSD
-F	--follow=name -retry	Same as <code>-f</code> + retry if file is rotated/truncated	GNU only; handles log rotation better
--retry		Keep trying to open file if inaccessible	GNU only
--pid=PID		Stop following when process PID dies	Useful for daemons
--max-unchanged-stats=N		(GNU) Check file attributes N times if no growth	Default 5

Classic live log monitoring:

Bash

```
tail -f /var/log/syslog      # watch system log in real time
tail -F /var/log/app.log     # survives log rotation (e.g., logrotate)
tail -f access.log | grep "ERROR" # live grep
```

With PID (stop when service dies):

Bash

```
tail -f --pid=$(pidof nginx) nginx.log
```

## 5. GNU Extensions (Linux) vs BSD (macOS/FreeBSD)

Important portability notes:

Feature	GNU (Linux)	BSD (macOS)	Workaround on macOS
-n -N (all but last N)	Yes (tail -n -5 = omit last 5)	No	head -n -5 (head has it)
-c +N / byte skipping	Yes	Limited	Use dd or sed
-F / --retry	Yes	No (-f only)	Install GNU tail via brew install coreutils → gtail
--max-unchanged-stats	Yes	No	N/A
Suffixes (K,M,G)	Yes	Yes (but case-sensitive)	

**macOS tip:** Prefix GNU versions with `g` after Homebrew install: `ghead`, `gtail`.

## 6. Real-World Examples

### 1. Quick file preview:

Bash

```
head large.csv          # first 10 rows
tail -n 20 huge.log     # last 20 log entries
```

### 1. Pagination / middle of file:

Bash

```
# Lines 100-110
sed -n '100,110p' file.txt
# Or combine head/tail
tail -n +100 file.txt | head -n 11
```

### 1. Binary files (e.g., images, executables):

Bash

```
head -c 16 /bin/ls          # first 16 bytes (magic number)
tail -c 100 binary.file     # last 100 bytes
```

### 1. Monitoring multiple logs:

Bash

```
tail -f log1.log log2.log      # interleaved output with headers
```

### 1. Skip header in CSV:

Bash

```
tail -n +2 data.csv | head -n 5  # rows 2-6 (skip header)
```

### 1. All but last 10 lines (GNU):

Bash

```
head -n -10 file.txt
# Or portable
head -n $(( $(wc -l < file.txt) - 10 )) file.txt
```

## 7. Common Pitfalls & Best Practices

- **Large files:** Both are extremely efficient — stream without loading full file.
- **Binary data:** Use `-c` for bytes; line mode may stop at first NUL byte.
- **Empty files:** Output nothing (graceful).
- **Trailing newline:** Preserved correctly.
- **Performance:** Faster than `sed` / `awk` for simple head/tail tasks.
- **Portability:** Avoid GNU-only features (`-F`, `-n` `-N`) in cross-platform scripts.
- **Live following ( -f ):** Consumes CPU if file doesn't change — combine with `timeout` or tools like `multitail`.
- **Multiple files:** Use `-v` for clarity.
- **Piping:** Perfect — e.g., `dmesg | tail -n 20` (last kernel messages).

### Best practice:

- Default 10 lines is often enough → just `head` file or `tail` file .
  - For logs → `tail -f` .
  - For skipping headers → `tail -n +2` .
  - For middle sections → `tail -n +START | head -n COUNT` .
- 

## 8. Quick Reference Table (All Major Options)

Option	Applies		
	To	Meaning	GNU/BSD
-n N / --lines=N	Both	N lines	Both
-n +N	tail	From line N	Both
-n -N	head/tail	All but last/first N (GNU)	GNU only
-c N / --bytes=N	Both	N bytes	Both
-c +N	tail	From byte N (GNU)	GNU only
-f / --follow	tail	Live follow	Both
-F	tail	Follow + retry (rotation)	GNU only
--pid=PID	tail	Stop when PID dies	Both
-q / --quiet	Both	No headers	Both
-v / --verbose	Both	Always headers	Both
-z / --zero-terminated	Both	NUL-terminated lines (rare)	GNU only

You've now got **complete mastery** of `head` and `tail` — from quick previews to live log monitoring.

What's your use case? Peeking at CSVs? Watching logs? Extracting middle sections?

Share a sample command or file type and I'll give you the perfect one-liner! 😊